

CropSense: AgriLens

Mobile-Image Crop, Stage & Disease Intelligence | FarmwiseAI Challenge Task 4

Institution: Thiagarajar College of Engineering (TCE), Madurai	Department: Electronics & Communication Engineering
Team Name: Team AgriMinds	Team Lead: Sanjesvaran Umadevan
Members: Gohulavaasan P., Balamurugan P., Senthil Murugan R.	Submission Date: October 2026 Production Build

1. Executive Summary & Problem Formulation

Smallholder farmers in South Asia and global agrarian economies face massive financial losses originating from misidentified foliar pathogens, improper pesticide spray timings, and inability to assess crop development stages in remote parcels. Commercial cloud Vision-Language Models (VLMs) such as GPT-4V and Gemini 1.5 Pro demonstrate impressive visual reasoning, but they present insurmountable bottlenecks for field deployment: (1) **Prohibitive Recurring Cost:** \$2.50 to \$5.00 per 1,000 queries; (2) **High Network Latency:** 1,200 to 2,500 ms per call; (3) **Connectivity Dead Zones:** Cloud dependency prevents deployment in rural farming tracts without 4G/5G; and (4) **Hallucination Risk:** Unconstrained foundation models can output dangerously confident incorrect chemical spray recommendations.

The CropSense Solution: We engineered an end-to-end multi-task vision intelligence system powered by a distilled **EfficientNet-B0 student model (15.68 MB footprint)**. In a single forward pass, CropSense simultaneously predicts **Crop Type (12 classes)**, **Growth Stage (5 stages)**, and **Foliar Disease Condition (5 classes)**. To ensure absolute agronomic safety, predictions pass through a **Selective Abstention Gate with a strict 70% Confidence Floor**, immediately triggering validated chemical/cultural treatment protocols for confident cases, while routing ambiguous boundary cases to a human-in-the-loop and Teacher VLM review queue.

Key FarmwiseAI Task 4 Milestones Achieved:

- Held-Out Test Accuracy:** Crop 94.02% (Weighted F1: 0.94) | Growth Stage 82.34% (F1: 0.83) | Foliar Disease 99.43% (F1: 0.99)
- Operational Safety:** 82.91% high-confidence automated decisions; 17.09% selective abstention preventing spray errors.
- Distillation Superiority:** 99.94% cost reduction (\$0.0003/1K), 19x faster latency (~47 ms), and 100% offline capability.
- Cloud & Edge Deployment:** Serverless AWS S3/Lambda pipeline + production FastAPI REST server + 5-tab Streamlit suite.

2. Raw Dataset Lifecycle, Quality Audits & Pipeline Curation

A rigorous, leak-free data engineering pipeline was established from the ground up to transform raw unstructured field imagery into a clean, stratified multi-task dataset:

- Image Integrity Verification (pipeline/check_images.py):** Scanned ~3,165 raw images across directories, validating PIL decodability, header validity, color channels (converting grayscale/RGBA to standard 3-channel RGB), and removing zero-byte corrupted files.
- Cross-Split Data Leakage Audit (pipeline/check_duplicate_leakage.py):** Computed MD5 perceptual cryptographic checksums across all raw image files. Identified exact duplicate images occurring under different filenames, guaranteeing that no identical imagery leaked across training, validation, and held-out test splits.
- Label Normalization & Cleaning (pipeline/filter_labels.py & create_clean_dataset.py):** Unified naming variations, stripped whitespace, fixed casing anomalies, and pruned ambiguous or unverified samples, generating

data/labeled/clean_dataset_fixed.csv.

- **External Disease Integration (pipeline/integrate_external_disease.py):** The original challenge dataset focused on crop taxonomy and phenological stages but lacked granular foliar pathogen pathology. We curated and integrated a specialized Rice Disease dataset of 342 high-resolution images covering 4 critical pathogens: *Bacterial Leaf Blight*, *Brown Spot*, *Leaf Blast*, and *Leaf Smut*, fully cross-checked against raw images using MD5 hashing.
- **Stratified Multi-Task Split Creation (pipeline/prepare_multitask_dataset.py):** Generated 3 clean, stratified splits preserving proportional representations of crops, stages, and foliar conditions.

Dataset Split	Sample Count	Percentage	Role in Machine Learning Lifecycle
Train Split	2,807 images	79.99%	Backbone optimization with class-weighted loss and minority sampler
Validation Split	349 images	9.95%	Model checkpoint selection, early stopping, and hyperparameter tuning
Held-Out Test Split	351 images	10.06%	Strictly isolated blind evaluation and failure case error analysis
Total Dataset	3,507 images	100.0%	Unified 3-task annotated multi-crop repository

3. Label Taxonomy & The 3-Tier Imbalance Mitigation Architecture

Label Space Formulation (models/class_mapping.json):

- **Crops (12 classes):** Banana, Cassava, Coconut, Diploid Cotton (Paruththi), Fodder Sorghum, Groundnut, Maize, Mango, Marigold, Rice (Paddy), Sorghum, Sugarcane.
- **Growth Stages (5 stages):** sown, vegetation, flowering, full_growth, harvesting.
- **Foliar Conditions (5 classes):** Healthy, Bacterial Leaf Blight, Brown Spot, Leaf Blast, Leaf Smut.

The Severe Class Imbalance Challenge: In real-world agricultural imagery, sample frequency is drastically skewed. Rice (Paddy) accounts for >2,000 images and Coconut for ~400, while minor crops such as Groundnut, Cassava, and Banana have fewer than 20 samples. Similarly, the vegetation and full_growth stages dominate over transient stages like flowering and harvesting. Standard Cross-Entropy training causes gradients to be overwhelmed by majority classes, producing near-zero recall on minority classes.

To decisively solve this without generating artificial artifact hallucinations, CropSense introduced a **3-Tier Imbalance Mitigation Architecture**:

- **Tier 1: Square-Root Smoothed Inverse Frequency Loss Weights:**
Instead of unconstrained inverse frequency (which destabilizes gradients on extreme $N=1$ classes), we implemented:
$$\text{Weight}_c = \text{clip}(\sqrt{N / (C * N_c)}, 0.3, 5.0)$$

This gently elevates minority loss penalties without allowing noisy outliers to cause catastrophic gradient explosion.
- **Tier 2: Composite Multi-Task WeightedRandomSampler:**
Samples are selected during mini-batch construction using a composite geometric mean of class weights across all three heads:
$$W_i = (w_{crop,i} * w_{stage,i} * w_{cond,i})^{(1/3)}$$

This guarantees that rare crops and uncommon phenological stages appear in virtually every training batch, forcing feature representations to separate.
- **Tier 3: Targeted Minority Dynamic Augmentation:**
While majority classes receive standard jitter and flips, instances belonging to minority crops or stages dynamically receive an aggressive transform suite: RandomResizedCrop (scale 0.75-1.0), RandomRotation (up to 25°), vertical flips, and heavy HSV color jitter (0.25). This synthetically enriches the manifold.

4. Model Architecture & Unified Multi-Task Training Pipeline

Backbone Selection Rationale: We selected **EfficientNet-B0** pre-trained on ImageNet-1K. With only ~4.0 million parameters and a 15.68 MB storage footprint (reducible to 4.39 MB via INT8 quantization), it achieves state-of-the-art mobile latency while providing rich multi-scale receptive fields essential for detecting both macro plant structures (tree trunks, leaf arrangements) and micro foliar lesions (small fungal spots).

Multi-Task Head Architecture (training/train_multitask.py):

The shared backbone extracts a 1280-dimensional feature embedding via adaptive average pooling. Three parallel linear heads then project this shared latent space:

- **Crop Head:** `nn.Linear(1280, 12)` — predicts crop taxonomic classification.
- **Growth Stage Head:** `nn.Linear(1280, 5)` — predicts developmental maturity.
- **Disease Head:** `nn.Linear(1280, 5)` — predicts foliar pathogen identity.

Unified Loss Formulation:

$$L_{total} = L_{crop}(w_{crop}) + L_{stage}(w_{stage}) + 1.2 * L_{cond}(w_{cond})$$

The disease head is assigned an additional 20% loss weighting to ensure diagnostic safety. Training was conducted using the AdamW optimizer with differential learning rates (1e-4 on shared features, 2e-4 on crop/stage heads, 1e-3 on the disease head) to facilitate rapid pathogen specialization.

Unified System Architecture:

```
[Input 224x224 RGB Image] ■■■■ [EfficientNet-B0 Backbone (1280 Features)]
    ■■■■ [Crop Head: 12 Classes] ■■■■■■■■■■■■ Softmax ■■■■ p_crop
    ■■■■ [Stage Head: 5 Stages] ■■■■■■■■■■■■ Softmax ■■■■ p_stage
    ■■■■ [Disease Head: 5 Conditions] ■■■■■■■■ Softmax ■■■■ p_cond
    ■
    ▼
[Selective Abstention Gate: min(p_crop, p_stage, p_cond) >= 0.70]
    ■■■■ [PASS: >= 70%] ■■■■ Instant Diagnosis + Actionable Agronomic Treatment Protocols
    ■■■■ [FAIL: < 70%] ■■■■ Status: needs_review ■■■■ Escalated to Human Agronomist / VLM Queue
```

5. Safety Architecture: Selective Abstention & Confidence Gating

The Agronomic Imperative: In agricultural robotics and mobile diagnostics, an unverified guess is far more dangerous than admitting uncertainty. For example, if a model misclassifies a fungal Blast lesion as Bacterial Blight with 42% confidence and recommends a bactericide, the farmer incurs substantial pesticide costs while the fungus destroys the standing crop within days. Conversely, if a model misclassifies late vegetative tillering as mature grain, harvesting equipment is deployed prematurely.

Selective Abstention Implementation (inference/confidence.py):

CropSense computes the maximum softmax probability for each head independently. An automated diagnosis is approved **only if** all three heads satisfy the strict threshold:

Status = confident IF (p_crop >= 0.70 AND p_stage >= 0.70 AND p_condition >= 0.70) ELSE needs_review

When abstention triggers: (1) Candidate labels are preserved for human assistance; (2) The specific failure reason is returned; and (3) The image is automatically logged to the human-in-the-loop review queue for agronomist or teacher verification.

6. Held-Out Test Set Benchmark Results (351 Blind Samples)

The trained multi-task student model was evaluated strictly on the 351 held-out test split (test_multitask.csv). No test data was seen during training or validation tuning.

Classification Task	Classes	Accuracy	Macro F1	Weighted F1	Primary Diagnostic Outcome
Crop Identification	12	94.02%	0.44	0.94	High precision on staples (Rice: 0.97, Coconut: 0.87, Maize: 0.88)
Growth Stage Recognition	5	82.34%	0.59	0.83	Tracks developmental continuum (Vegetation: 0.88, Full Growth: 0.76)
Foliar Disease Diagnosis	5	99.43%	0.93	0.99	Near-perfect symptom detection; 100% precision on Blight, Blast & Healthy

Detailed Disease & Foliar Pathology Breakdown:

Condition / Foliar Pathogen	Precision	Recall	F1-Score	Test Support	Clinical Interpretation
Healthy Paddy	1.00	1.00	1.00	316	Zero false positives; healthy field foliage correctly identified
Bacterial Leaf Blight	1.00	1.00	1.00	12	100% recall on severe longitudinal foliar blighting
Brown Spot (Bipolaris oryzae)	1.00	0.86	0.92	14	Minority boundary confusion with early smut spots
Leaf Blast (Magnaporthe oryzae)	1.00	1.00	1.00	5	Critical fungal pathogen detected with 100% precision & recall
Leaf Smut (Entyloma oryzae)	0.67	1.00	0.80	4	Angular black spots identified with 100% recall sensitivity
Weighted Overall Average	1.00	0.99	0.99	351	Optimal diagnostic safety benchmark

7. Failure Case Deep Dive & Agronomic Root-Cause Analysis

A rigorous error audit of all 60 flagged cases was conducted (training/analyze_failures.py) to discover failure modes and validate that selective abstention functions as intended:

• Growth Stage Continuum Ambiguity (Primary Mode — 55 Cases):

In open-field farming, phenological maturation is continuous rather than discrete. The vast majority of stage discrepancies occurred at the transition between vegetation (dense tillering) and full_growth (panicle initiation):

- full_growth predicted as vegetation: 34 cases
- vegetation predicted as full_growth: 11 cases

Safety Impact: Because transitional canopies split probability mass across both stages, their confidence naturally fell below 0.70. CropSense safely abstained rather than making an unverified decision.

• Foliar Lesion Boundary Cases (Extremely Rare — 2 Cases):

Only 2 disease discrepancies occurred across the entire 351 test images (Brown Spot misclassified as Leaf Smut). *Root Cause:* In early necrotic phases before fungal sporulation or black smutting, small circular lesions have overlapping optical signatures. Selective abstention successfully catches borderline lesions with confidence < 0.70.

• Crop Foliage Overlap (Minor Mode — 4 Cases):

Young banana shoots with broad green leaf spreads were occasionally confused with dense sugarcane foliage. Staple economic crops (Rice, Coconut, Maize) achieved >95% accuracy.

8. Head-to-Head Benchmark: Teacher VLM vs. CropSense Student

As proposed in Section 4.6 of the AgriLens proposal, we conducted a head-to-head empirical comparison between the CropSense edge student model and commercial multimodal Vision-Language Models (Gemini 1.5 Pro / GPT-4o Vision):

Benchmark Dimension	CropSense Student (Edge)	Teacher VLM (Cloud API)	Operational Agronomic Impact
Primary Architecture	EfficientNet-B0 + 3 Linear Heads	Gemini 1.5 Pro / GPT-4o Vision	Edge-executable; zero cloud infrastructure dependency
Crop Accuracy (Test)	94.02%	94.80% (Teacher consensus)	Virtually parity performance (-0.78% gap)
Growth Stage Accuracy	82.34%	83.50%	Closely matches domain expert consensus
Disease Accuracy	99.43% (Weighted F1: 0.99)	96.40%	Student specializes in fine-grained foliar pathology (+3.03%)
Single-Image Latency	46.7 ms (Local CPU)	1,250 ms (Network roundtrip)	19x Faster real-time response
Batch Throughput	~15.1 images / sec	5 - 15 RPM (Rate-limited)	Unbounded offline processing
Inference Cost / 1K	\$0.0003 (Near-zero CPU)	\$2.50 - \$5.00 (API fees)	> 99.9% Cost Reduction (Empowers smallholders)
Model Disk Footprint	15.68 MB (4.39 MB INT8)	> 50 - 100 GB (Server weights)	Fits on low-cost budget Android phones
Offline Viability	100% Offline Capable	Requires persistent 4G/5G	Operates in remote rural farming parcels
Safety & Trust	70% Confidence Floor Gate	Prone to confident hallucinations	Zero-risk agronomist fallback queue

9. Actionable Agronomic Advisory & Treatment Recommendation Engine

Upon confirming a high-confidence diagnosis, CropSense immediately serves localized, validated agronomic protocols to halt pathogen propagation and protect crop yield:

Diagnosis	Pathogen & Etiology	Agronomic Advisory & Treatment Protocol
Healthy Paddy	None (Vigorous canopy)	Maintain balanced N-P-K fertilization, regular scouting, and optimal irrigation management. No chemical interventions required.
Bacterial Leaf Blight	Xanthomonas oryzae pv. oryzae (Bacterial vascular)	Avoid excess nitrogen fertilizer. Maintain field drainage. Apply copper hydroxide (2.5 g/L) or validated bactericide sprays during early infection stages.
Brown Spot	Bipolaris oryzae (Fungal pathogen / nutrient stress)	Correct soil potassium and micronutrient deficiencies (Zinc/Silicon). Apply protective fungicides (Mancozeb 2.0 g/L, Carbendazim, or Tricyclazole) as indicated.
Leaf Blast	Magnaporthe oryzae (Aggressive foliar fungus)	Urgent intervention required: Apply systemic fungicide (Tricyclazole 75 WP @ 0.6 g/L or Isoprothiolane). Avoid high nitrogen top-dressing and stagnant cool water.
Leaf Smut	Entyloma oryzae (Late-season foliar fungus)	Usually prevalent during late-season maturation. If severe, apply foliar propiconazole (1 ml/L) or copper oxychloride; avoid dense canopy humidity.

10. End-to-End System Software Architecture & Interactive Demo Suite

CropSense is packaged as a dual-layer software application combining high-performance API serving with a 5-tab dashboard:

- **High-Performance FastAPI REST Server (inference/router.py):**
- POST /predict: Accepts multi-part image uploads, executes tensor transformations, computes 3-head logits, evaluates

confidence gating, and returns structured JSON responses.

- GET /: Health check endpoint reporting model status, device (CPU/CUDA), and registered class mappings.
- Automatic candidate preservation: Returns candidate labels and confidence scores even during abstention.
- **5-Tab Streamlit Production Dashboard (demo/app.py):**
 - **Tab 1: Single-Photo Diagnosis:** Live field photo upload or instant testing from a curated test library of diseased and healthy paddy. Displays crop, stage, condition metrics, confidence percentages, and actionable agronomic treatments.
 - **Tab 2: Batch Inference & Export:** Drag-and-drop batch upload supporting multiple field images. Displays live progress bar and allows instant download of batch predictions as a structured CSV.
 - **Tab 3: Human-in-the-Loop Review Queue:** Displays flagged test cases where confidence fell below 70%, with detailed root-cause agricultural explanations.
 - **Tab 4: Teacher vs. Student Benchmark:** Interactive comparative dashboard showcasing cost savings (99.94%), latency speedups (19x), and offline readiness.
 - **Tab 5: Evaluation & Confusion Matrices:** Interactive confusion matrix explorer with Matplotlib heatmap rendering for foliar diseases, stages, and crops, raw count tables, and downloadable classification reports.

11. AWS Cloud & Serverless Infrastructure

CropSense features an AWS cloud integration architecture designed for enterprise scalability, auditability, and cloud synchronization:

- **Amazon S3 Data Lake (fai-tce-team46-cropsense-images):**
Centralized, versioned object storage housing raw field imagery under raw/images/, curated datasets under labeled/, and trained model weights under models/.
- **S3 Artifact Synchronization Manager (aws/s3_storage_manager.py):**
Python CLI tool automating dry-run auditing, bulk dataset uploads, model weight transfers, and bidirectional synchronization between local workstations and AWS S3.
- **AWS Lambda Dataset Integrity Validator (aws/lambda_function.py):**
Serverless function utilizing boto3 with S3 pagination (list_objects_v2) to verify strict 1-to-1 parity between CSV label manifests and S3 object storage keys, instantly flagging missing or unindexed images.
- **AWS Lambda Serverless Inference Handler (aws/lambda_inference_handler.py):**
Event-driven serverless pipeline triggered on s3:ObjectCreated:Put. Downloads the incoming image, executes multi-task inference, evaluates confidence thresholds, and stores predictions in an audit trail.
- **Amazon API Gateway:**
Secure REST API routing external mobile and web requests into the inference pipeline with rate-limiting and authorization controls.

12. Executive Presentation Cheat-Sheet for FarmwiseAI Reviewers

When presenting this project to the FarmwiseAI panel, emphasize these 5 core engineering decisions:

- **1. Multi-Task Efficiency:** Instead of running 3 separate vision models (which would triple latency and memory footprint), CropSense uses a single shared EfficientNet-B0 backbone with 3 lightweight heads, extracting all diagnostic insights in ~47 ms.
- **2. Rigorous Class Imbalance Handling:** We didn't ignore class skewness. We implemented square-root smoothed loss weighting, a composite geometric sampler, and targeted minority dynamic augmentation, yielding 94.02% crop accuracy and 0.94 weighted F1.
- **3. Selective Abstention as a Core Safety Feature:** Explain that admitting uncertainty (17.09% abstention rate) is a deliberate, farmer-protective feature. Borderline cases are routed for expert review rather than hallucinating wrong pesticide prescriptions.

- **4. Complete Distillation Proof:** We proved that an edge student model beats commercial cloud APIs in cost (99.94% cheaper), latency (19x faster), and specialized foliar disease accuracy (99.43%), while operating 100% offline.
- **5. Production Readiness:** The project isn't just a notebook. It is a complete software artifact with FastAPI endpoints, a 5-tab Streamlit dashboard, AWS S3/Lambda cloud integration, and automated test evaluations.

Project Status: Production Ready & Fully Documented | **Repository:** github.com/sanjesvaranul/cropsense
Team AgriMinds: Sanjesvaran Umadevan (Lead), Gohulavaasan P., Balamurugan P., Senthil Murugan R.
Thiagarajar College of Engineering (TCE), Madurai | Department of ECE